

Lab 4 — create simple container

Intro

- The idea of the work is to `clone()` process with flags enabling the separate namespaces for it, etc., to prepare rootfs image for this process to `chroot` into, to configure `cgroups`.

Features

- Rootfs is based on Ubuntu 20.04 base image. Sysbench is added to the container filesystem image through `init_container.sh`. On startup, `bash` shell is invoked.
- Container is created with its own namespaces:
 1. PID namespace (`CLONE_NEWPID`): The new process will have its own PID namespace. Processes in this namespace can only see the processes within the same namespace. The first process in this namespace is usually the init process, with a PID of 1.
 2. UTS namespace (`CLONE_NEWUTS`): The new process will have its own UTS namespace, which includes the hostname and domain name. This allows a process to have a different hostname inside and outside the namespace.
 3. Network namespace (`CLONE_NEWNET`): The new process will have its own network namespace. We create namespace within Bash script using `ip netns add` and further set it up. This means that it will have its own set of network interfaces, IP addresses, routing tables, and firewall rules, independent of the host and other processes. A pair of virtual interfaces is created: `veth_host` and `veth_container`, and assigned IP addresses 192.168.10.1 and 192.168.10.2, accordingly. The `container` binary simply joins existing namespace `container_network_ns`.
 4. Mount namespace (`CLONE_NEWNS`): The new process will have its own mount namespace. This means that it will have its own filesystem root directory and its own set of mount points, independent of the host and other processes.

```

artem@latitude:~/OneDrive/Inno/S23/TVL/container_lab (master) $ sudo ./build/container
[sudo] password for artem:
root@container:/# pwd
/
root@container:/# ps
  PID TTY          TIME CMD
    1 ?           00:00:00 container
    2 ?           00:00:00 bash
    5 ?           00:00:00 ps
root@container:/# ls
bin  dev  home  lib32  libx32  media  opt  root  sbin  sys  usr
boot  etc  lib  lib64  lost+found  mnt  proc  run  srv  tmp  var
root@container:/#

```

Filesystem isolation

```

/ # ls
bin      lib      proc     sbin     usr
dev      lost+found  root     sys      var
etc      opt      run      tmp

/ # ls
bin      lib      proc     sbin     usr
dev      lost+found  root     sys      var
etc      opt      run      tmp

/ # ls
bin      lib      proc     sbin     usr
dev      lost+found  root     sys      var
etc      opt      run      tmp

/ # touch nohost.txt
/ # ls
bin      lib      opt      run      tmp
dev      lost+found  proc     sbin     usr
etc      nohost.txt  root     sys      var

/ # ls
bin      lib      opt      run      tmp
dev      lost+found  proc     sbin     usr
etc      nohost.txt  root     sys      var
/ # █

```

```

root@latitude:/home/artem $ cd /
root@latitude:/ $ ls
bin  home  lib32  media  root  srv  var
boot  initrd.img  lib64  mnt  run  sys  vmlinuz
dev  initrd.img.old  libx32  opt  sbin  tmp  vmlinuz.old
etc  lib  lost+found  proc  snap  usr
root@latitude:/ $ touch nocontainer.txt
root@latitude:/ $ ls
bin  home  lib32  media  proc  snap  usr
boot  initrd.img  lib64  mnt  root  srv  var
dev  initrd.img.old  libx32  nocontainer.txt  run  sys  vmlinuz
etc  lib  lost+found  opt  sbin  tmp  vmlinuz.old
root@latitude:/ $ ls
bin  home  lib32  media  proc  snap  usr
boot  initrd.img  lib64  mnt  root  srv  var
dev  initrd.img.old  libx32  nocontainer.txt  run  sys  vmlinuz
etc  lib  lost+found  opt  sbin  tmp  vmlinuz.old
root@latitude:/ $

```

PID isolation

```

/ # ps aux
PID  USER  TIME  COMMAND
  1  root    0:00  cmake-build-debug/container
  2  root    0:00  /bin/sh
 11  root    0:00  ps aux
/ # █

```

```

root@latitude:/ $ ps aux
USER      PID  %CPU  %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
root         1  0.0  0.0 166032 11356 ?        Ss   May17   0:05 /sbin/
root         2  0.0  0.0      0   0 ?        S    May17   0:00 [kthre
root         3  0.0  0.0      0   0 ?        I<   May17   0:00 [rcu_g
root         4  0.0  0.0      0   0 ?        I<   May17   0:00 [rcu_p
root         6  0.0  0.0      0   0 ?        I<   May17   0:00 [kwork
root         8  0.0  0.0      0   0 ?        I<   May17   0:00 [mm_pe
root         9  0.0  0.0      0   0 ?        S    May17   0:00 [rcu_t
root        10  0.0  0.0      0   0 ?        S    May17   0:00 [rcu_t
root        11  0.0  0.0      0   0 ?        S    May17   0:01 [ksoft
root        12  0.0  0.0      0   0 ?        I    May17   0:41 [rcu_s
root        13  0.0  0.0      0   0 ?        S    May17   0:00 [migr
root        15  0.0  0.0      0   0 ?        S    May17   0:00 [cpuhp

```

Network isolation (and communication)

```

/ # ip a
1: lo: <LOOPBACK> mtu 65536 qdisc noop state DOWN qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
37: veth_container@if38: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 q
disc noqueue state UP qlen 1000
    link/ether e6:dd:e0:f3:5a:5b brd ff:ff:ff:ff:ff:ff
    inet 192.168.10.2/24 scope global veth_container
        valid_lft forever preferred_lft forever
    inet6 fe80::e4dd:e0ff:fef3:5a5b/64 scope link
        valid_lft forever preferred_lft forever
/ # ping 192.168.10.1 -c 1
PING 192.168.10.1 (192.168.10.1): 56 data bytes
64 bytes from 192.168.10.1: seq=0 ttl=64 time=0.208 ms

--- 192.168.10.1 ping statistics ---
1 packets transmitted, 1 packets received, 0% packet loss
round-trip min/avg/max = 0.208/0.208/0.208 ms
/ # []

30: vpn: <POINTOPOINT,NOARP,UP,LOWER_UP> mtu 1280 qdisc noqueue state UNKNOW
N group default qlen 1000
    link/none
    inet 10.242.104.10/24 scope global vpn
        valid_lft forever preferred_lft forever
36: vethdb728cf@if35: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noque
ue master docker0 state UP group default
    link/ether de:f2:b7:a3:25:b9 brd ff:ff:ff:ff:ff:ff link-netnsid 1
38: veth_host@if37: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue
state UP group default qlen 1000
    link/ether 92:74:1a:b1:95:71 brd ff:ff:ff:ff:ff:ff link-netns container_
network_ns
    inet 192.168.10.1/24 scope global veth_host
        valid_lft forever preferred_lft forever
root@latitude:/ # ping 192.168.10.2 -c 1
PING 192.168.10.2 (192.168.10.2) 56(84) bytes of data.
64 bytes from 192.168.10.2: icmp_seq=1 ttl=64 time=0.104 ms

--- 192.168.10.2 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 0.104/0.104/0.104/0.000 ms

```

```

container:/# telnet 192.168.10.1 2222
Trying 192.168.10.1...
Connected to 192.168.10.1.
Escape character is '^]'.
hello from container
[]

8: docker0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state U
P group default
    link/ether 02:42:67:0b:e0:c8 brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.1/16 brd 172.17.255.255 scope global docker0
        valid_lft forever preferred_lft forever
30: vpn: <POINTOPOINT,NOARP,UP,LOWER_UP> mtu 1280 qdisc noqueue state UNKNOW
N group default qlen 1000
    link/none
    inet 10.242.104.10/24 scope global vpn
        valid_lft forever preferred_lft forever
36: vethdb728cf@if35: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noque
ue master docker0 state UP group default
    link/ether de:f2:b7:a3:25:b9 brd ff:ff:ff:ff:ff:ff link-netnsid 1
42: veth_host@if41: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue
state UP group default qlen 1000
    link/ether 92:74:1a:b1:95:71 brd ff:ff:ff:ff:ff:ff link-netns container_
network_ns
    inet 192.168.10.1/24 scope global veth_host
        valid_lft forever preferred_lft forever
root@latitude:/ # nc -l -s 192.168.10.1 -p 2222
hello from container
[]

```

Comparison with Docker

- CPU info:

```

artem@latitude:~/OneDrive/Inno/S23/TVL/container_lab (master) $ sudo ./cmake-build-debug/container
[sudo] password for artem:
starting container
container:/# cat /proc/cpuinfo
processor       : 0
vendor_id     : GenuineIntel
cpu family    : 6
model         : 61
model name    : Intel(R) Core(TM) i5-5300U CPU @ 2.30GHz
stepping      : 4
microcode    : 0x2f
cpu MHz      : 2175.051
cache size   : 3072 KB
physical id   : 0
siblings     : 4
core id      : 0
cpu cores    : 2
apicid       : 0
initial apicid : 0
fpu          : yes
fpu_exception : yes
cpuid level  : 20
wp           : yes
flags        : fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush
dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx pdpe1gb rdtscp lm constant_tsc arch_perfmon p
s bts rep_good nopl xtopology nonstop_tsc cpuid aperfperf pni pclmulqdq dtes64 monitor ds cpl vmx
smx est tm2 ssse3 sbdb fma cx16 xtpr pdcm pcid sse4_1 sse4_2 x2apic movbe popcnt tsc_deadline_t
im er aes xsave avx f16c rdrand lahf_lm abm 3dnowprefetch cpuid_fault epb invpcid_single ssbd ibrs
ibpb stibp tpr_shadow vnmi flexpriority ept vpid ept_ad fsgsbase tsc_adjust bmi1 hle avx2 smep b
mi2 erms invpcid rtm rdseed adx smap intel_pt xsaveopt dtherm ida arat pln pts md_clear flush_l
id vmx flags : vnmi preemption_timer invpcid ept_x_only ept_ad ept_1gb flexpriority tsc_offset
vtr mtrf vapid ept vpid unrestricted_guest ple shadow_vmcs
bugs          : cpu_meltdown spectre_v1 spectre_v2 spec_store_bypass l1tf mds swapgs taa itlb_mu
lthit srbsds mmio_unknow
artem@latitude:~$ docker run -it ubuntu sh
#
# cat /proc/cpuinfo
processor       : 0
vendor_id     : GenuineIntel
cpu family    : 6
model         : 61
model name    : Intel(R) Core(TM) i5-5300U CPU @ 2.30GHz
stepping      : 4
microcode    : 0x2f
cpu MHz      : 2194.951
cache size   : 3072 KB
physical id   : 0
siblings     : 4
core id      : 0
cpu cores    : 2
apicid       : 0
initial apicid : 0
fpu          : yes
fpu_exception : yes
cpuid level  : 20
wp           : yes
flags        : fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush
dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx pdpe1gb rdtscp lm constant_tsc arch_perfmon p
s bts rep_good nopl xtopology nonstop_tsc cpuid aperfperf pni pclmulqdq dtes64 monitor ds cpl vmx
smx est tm2 ssse3 sbdb fma cx16 xtpr pdcm pcid sse4_1 sse4_2 x2apic movbe popcnt tsc_deadline_t
im er aes xsave avx f16c rdrand lahf_lm abm 3dnowprefetch cpuid_fault epb invpcid_single ssbd ibrs
ibpb stibp tpr_shadow vnmi flexpriority ept vpid ept_ad fsgsbase tsc_adjust bmi1 hle avx2 smep b
mi2 erms invpcid rtm rdseed adx smap intel_pt xsaveopt dtherm ida arat pln pts md_clear flush_l
id vmx flags : vnmi preemption_timer invpcid ept_x_only ept_ad ept_1gb flexpriority tsc_offset
vtr mtrf vapid ept vpid unrestricted_guest ple shadow_vmcs
bugs          : cpu_meltdown spectre_v1 spectre_v2 spec_store_bypass l1tf mds swapgs taa itlb_mu
lthit srbsds mmio_unknow

```

- Networking info:

```

artem@latitude:~/OneDrive/Inno/S23/TVL/container_lab (master) $ sudo ./cmake-build-debug/container
starting container
container:/# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state DOWN qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
4: veth_container@if5: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 qdisc noqueue state UP ql
    en 1000
    link/ether e6:dd:e0:f3:5a:5b brd ff:ff:ff:ff:ff:ff
    inet 192.168.10.2/24 scope global veth_container
        valid_lft forever preferred_lft forever
    inet6 fe80::e4dd:e0ff:fe3:5a5b/64 scope link
        valid_lft forever preferred_lft forever
container:/#

```

```

artem@latitude:~$ docker run -it alpine sh
/# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
17: eth@if18: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 qdisc noqueue state UP
    link/ether 02:42:ac:11:00:02 brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.2/16 brd 172.17.255.255 scope global eth0
        valid_lft forever preferred_lft forever
/#

```

Benchmark

My container	Ubuntu Docker	Host
<pre> root@container:/# sysbench cpu --cpu-max-prime=20000 run sysbench 1.1.0-4228c85 (using bundled LuaJIT 2.1.0-beta3) Running the test with following options: Number of threads: 1 Initializing random number generator from current time Prime numbers limit: 20000 Initializing worker threads... Threads started! CPU speed: events per second: 270.65 Throughput: events/s (eps): 270.6543 time elapsed: 10.0017s total number of events: 2707 Latency (ms): min: 3.56 avg: 3.69 max: 16.29 95th percentile: 3.89 sum: 9998.66 Threads fairness: events (avg/stddev): 2707.0000/0.00 execution time (avg/stddev): 9.9987/0.00 </pre>	<pre> root@b0e0ecb345b1:/mnt/shared# src/sysbench cpu --cpu-max-prime=20000 run sysbench 1.1.0-4228c85 (using bundled LuaJIT 2.1.0-beta3) Running the test with following options: Number of threads: 1 Initializing random number generator from current time Prime numbers limit: 20000 Initializing worker threads... Threads started! CPU speed: events per second: 265.57 Throughput: events/s (eps): 265.5731 time elapsed: 10.0010s total number of events: 2656 Latency (ms): min: 3.55 avg: 3.76 max: 16.45 95th percentile: 4.33 sum: 9996.82 Threads fairness: events (avg/stddev): 2656.0000/0.00 execution time (avg/stddev): 9.9968/0.00 </pre>	<pre> artem@latitude:~/OneDrive/Inno/S23/TVL/container_lab (master) \$ sysben ch cpu --cpu-max-prime=20000 run sysbench 1.1.0-4228c85 (using bundled LuaJIT 2.1.0-beta3) Running the test with following options: Number of threads: 1 Initializing random number generator from current time Prime numbers limit: 20000 Initializing worker threads... Threads started! CPU speed: events per second: 278.07 Throughput: events/s (eps): 278.0728 time elapsed: 10.0010s total number of events: 2701 Latency (ms): min: 3.56 avg: 3.78 max: 17.45 95th percentile: 3.89 sum: 9998.16 Threads fairness: events (avg/stddev): 2701.0000/0.00 execution time (avg/stddev): 9.9982/0.00 </pre>

Table with metrics

	command executed	my container	Docker (ubuntu 22.04)	host machine
CPU total time	<code>sysbench cpu --cpu-max-prime=20000 run</code>	9.9899 s	9.9968 s	9.9982 s
File IO write	<code>sysbench fileio --file-total-size=1G --file-num=128 --file-test-mode=seqwr run</code>	1073741824 bytes written in 6.60 seconds (155.11 MiB/sec).	1073741824 bytes written in 5.06 seconds (202.51 MiB/sec).	1073741824 bytes written in 7.89 seconds (129.82 MiB/sec).
File IO read	<code>sysbench fileio --file-total-size=1G --file-num=128 --file-test-mode=seqrd run</code>	IOPS=308201.48 4815.65 MiB/s (5049.57 MB/s)	IOPS=323803.01 5059.42 MiB/s (5305.19 MB/s)	IOPS=324396.52 5068.70 MiB/s (5314.91 MB/s)
Memory access	<code>sysbench memory --memory-block-size=1K --memory-total-size=4G run</code>	0.4148 s	0.4163 s	0.4164

- After conducting several tests, it was found that the performance of the container created in this lab and Docker's Ubuntu 22.04 image differed insignificantly. The CPU time, file IO write and read, and memory access of both containers were similar. The reason for this is that mechanism I used in my container are nearly identical to those used in Docker. Overall, the container in this lab and Docker's Ubuntu 22.04 image performed similarly, with the slight edge going to Docker due to its more optimized storage driver.
- One of the difficulties that worth highlighting in this lab was figuring out how to set up the network within the container. It was also found that the capabilities required to run containers are quite high, and the simplest way to run a container is to run as root. Creating an appropriate rootfs with all necessary utilities was also a challenge, but a script provided by Alpine Linux was used to create a rootfs with sysbench, telnet, and ping: <https://github.com/alpinelinux/alpine-make-rootfs>. This rootfs is stored as a dependency in `deps/` directory of the GitHub repo.

Links

- This project on Github: <https://github.com/ar7ch/lab4tv>

Sources

1. <https://man7.org/linux/man-pages/man7/namespaces.7.html>
2. <https://cesarvr.io/post/2018-05-22-create-containers/>
3. <https://github.com/akopytov/sysbench#general-syntax>
4. <https://docs.docker.com/storage/storagedriver/>
5. <https://man7.org/linux/man-pages/man8/ip-netns.8.html>